

Anatomy of a Real-Time Elixir App

Stephen Bussey

Hi, I'm Steve Bussey

- Software Architect at SalesLoft (Atlanta, GA, USA)
- Elixir in production for several years
- I love Elixir!



The
Pragmatic
Programmers

Real-Time Phoenix

Build Highly Scalable Systems
with Channels



Stephen Bussey

Series editor: *Bruce A. Tate*

Development editor: *Jacquelyn Carter*

<https://sal.es/rtp>

Goals for Today

- Understand importance of reading the source code of libraries that we use
- Dive deep into the basics of how Phoenix Channels work

Real-Time Elixir Apps

What is a Real-Time App?

- Interact with clients and keep them up to date as data changes
- Goals set around how long it takes for data changes to propagate from source to client
- A real-time app should ideally always reflect the “truth”

Components of a Real-Time App

- Client (JavaScript, iPhone app, etc)
- Server-Client connection (ex. WebSocket)
- Server (Elixir, Phoenix)
- Data Pipeline (GenStage, homemade)

We're going to look at and break down the server component today. It would take awhile to do this for all of the different components!

Anatomy

What is anatomy?

— — —

- “A study of the structure or internal workings of something.”
- Concerned with the relationship between small parts of the whole

Anatomy of a Software Library

— — —

- We can get a clearer picture of the whole by understanding the parts
- We can more easily dig in and debug problems that pop up
- We gain self-sufficiency in our usage of a library

Self-sufficiency takes us to the next level

Self-sufficient developers

— — —

- Can answer questions that may not be immediately clear
- Can more easily contribute to open-source or internal codebases
- Get to the bottom of problems (bugs) to find a path forward

Our Path to Self-Sufficiency

— — —

- Approach a library with curiosity
- Dig in when a problem appears
- Contribute back any learnings, such as documentation or found defects

Phoenix Channels

Implementing the 101 of Channels

- Basic Socket
- Channel with a message handler
- JavaScript connects to the Socket and Channel
- Let's view that code

Lots of moving parts, not much code

- What is a Socket, and how does it connect?
- How does a Channel differ from a Socket?
- How does a Channel become joined?
- How does a message get processed and responded to?

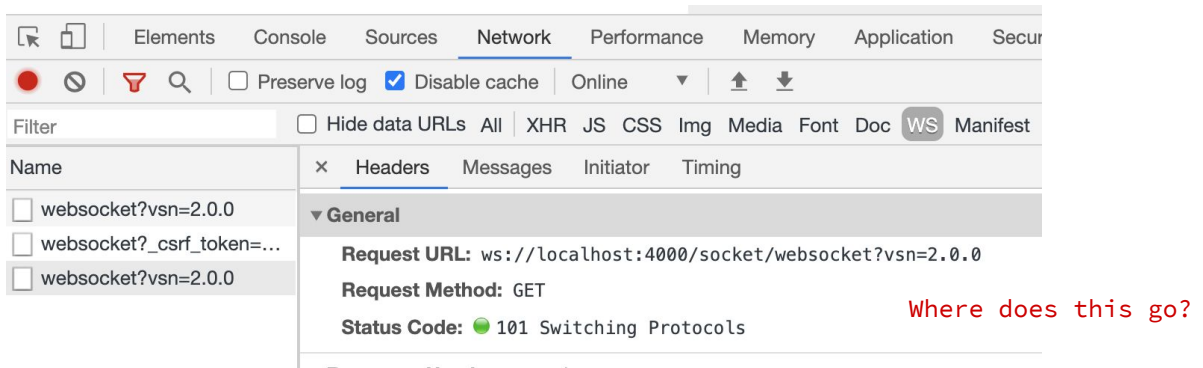
Let's dig into the code

- We'll only go deeper if we have a clear reason to (no magic jumps)
- Assume no prior knowledge, we have to see it ourselves
- We may end up with more questions than answers, but we'll answer our key questions

Socket connection and routing

WebSocket is used to connect to the Socket

— — —



Socket Definition

- Defined /socket in **DemoWeb.Endpoint**
- Leverages the socket/3 function of **Phoenix.Endpoint**

```
socket "/socket", DemoWeb.UserSocket,  
  websocket: true,  
  longpoll: false
```

Phoenix.Endpoint defmacro socket/3

```
848   defmacro socket(path, module, opts \\ []) do
849     # Tear the alias to simply store the root in the AST.
850     # This will make Elixir unable to track the dependency
851     # between endpoint <=> socket and avoid recompiling the
852     # endpoint (alongside the whole project ) whenever the
853     # socket changes.
854     module = tear_alias(module)
855
856     quote do
857       • @phoenix_sockets {unquote(path), unquote(module), unquote(opts)}
858     end
859   end
```

How is phoenix_sockets used?

Phoenix.Endpoint before compile

```
585   @doc false
586   defmacro __before_compile__({module: module}) do
587     sockets = Module.get_attribute(module, :phoenix_sockets)
588
589     dispatches =
590       for {path, socket, socket_opts} <- sockets,
591         {path, type, conn_ast, socket, opts} <- socket_paths(module, path, socket, socket_opts) do
592         quote do
593           defp do_handler(unquote(path), conn, _opts) do
594             {unquote(type), unquote(conn_ast), unquote(socket), unquote(Macro.escape(opts))}
595           end
596         end
597       end
598   end
```

Each dispatch becomes a quoted function

Phoenix.Endpoint dispatches

```
624     @doc false
625     def __handler__({path_info: path} = conn, opts), do: do_handler(path, conn, opts)
626     unquote(dispatches)
627     defp do_handler(_path, conn, opts), do: {:plug, conn, __MODULE__, opts}

        Pretty cool
```

When Endpoint is called, it will execute the dispatch functions, instead of going into Plug

Phoenix.Endpoint.socket_paths

```
---
631 defp socket_paths(endpoint, path, socket, opts) do
632   paths = []
633   websocket = Keyword.get(opts, :websocket, true)
634   longpoll = Keyword.get(opts, :longpoll, false)
635
636   paths =
637     if websocket do
638       config = Phoenix.Socket.Transport.load_config(websocket, Phoenix.Transports.WebSocket)
639       {conn_ast, match_path} = socket_path(path, config)
640       [{match_path, :websocket, conn_ast, socket, config} | paths]
641     else
642       paths
643     end
644
645   paths =
646     if longpoll do
647       config = Phoenix.Socket.Transport.load_config(longpoll, Phoenix.Transports.LongPoll)
```

A WebSocket is configured via `Transports.WebSocket`, which contributes to building up the dispatch AST

Output of do_handler

- Inspected by modifying dependency source code

What the heck? No function calls

```
{:defp, [context: Phoenix.Endpoint, import: Kernel],
[
  {:do_handler, [context: Phoenix.Endpoint],
  [
    ["socket", "websocket"],
    {:conn, [], Phoenix.Endpoint},
    {:_opts, [], Phoenix.Endpoint}
  ]},
  [
    do: {:{}, [],
    [
      :websocket,
      {:conn, [], Phoenix.Endpoint},
      DemoWeb.UserSocket,
      [
        path: "/websocket",
        serializer: [
          {Phoenix.Socket.V1.JSONSerializer, "~> 1.0.0"},
          {Phoenix.Socket.V2.JSONSerializer, "~> 2.0.0"}
        ],
        timeout: 60000,
        transport_log: false,
        compress: false
      ]
    ]
  ]}]
}]}
```

This is what a handler AST looks like

Coming Together

- Neat metaprogramming techniques!
- Phoenix maps each defined socket into a set of dispatch functions
- Phoenix supports WebSocket and LongPoll
- But we're seemingly stuck

Grep is curiosity's best friend

- `__handler__` is very specific
- Only appears in `endpoint.ex` and a `cowboy2_handler.ex`

```
→ demo grep -R "__handler__" deps/  
deps//phoenix/lib/phoenix/endpoint/cowboy2_handler.ex:      case endpoint.__handler__(conn, opts) do  
deps//phoenix/lib/phoenix/endpoint/cowboy2_handler.ex:      [{}^endpoint, :__handler__, _, _] | _] when reason == :undef and retry? ->  
deps//phoenix/lib/phoenix/endpoint.ex:      def __handler__({path_info: path} = conn, opts), do: do_handler(path, conn, opts)  
→ demo █
```

Now we're unstuck

[Cowboy2Handler init/4](#)

Transports.Websocket controls connections

cowboy_websocket is the response

```
18 defp init(conn, endpoint, opts, retry?) do
19   try do
20     case endpoint.__handler__(conn, opts) do
21       {:websocket, conn, handler, opts} ->
22         case Phoenix.Transports.WebSocket.connect(conn, endpoint, handler, opts) do
23           {:ok, %Plug.Conn{adapter: {@connection, req}} = conn, state} ->
24             cowboy_opts =
25               opts
26               |> Enum.flat_map(fn
27                 {timeout, timeout} -> [idle_timeout: timeout]
28                 {:compress, _} = opt -> [opt]
29                 {:max_frame_size, _} = opt -> [opt]
30                 _other -> []
31               end)
32               |> Map.new()
33
34             {:cowboy_websocket, copy_resp_headers(conn, req), [handler | state], cowboy_opts}
35
36           {:error, %Plug.Conn{adapter: {@connection, req}} = conn} ->
37             {:ok, copy_resp_headers(conn, req), {handler, opts}}
38         end
39     end
40   end
41 end
```

Each WebSocket is processed by the Transport.connect function and returned as a cowboy websocket

Fork in the road. We need to know about the Transport and also Cowboy

Transports.WebSocket

Hey, it's our Socket!
Sort of...

```
16 def connect(%{method: "GET"} = conn, endpoint, handler, opts) do
17   conn
18   |> Plug.Conn.fetch_query_params()
19   |> Transport.code_reload(endpoint, opts)
20   |> Transport.transport_log(opts[:transport_log])
21   |> Transport.force_ssl(handler, endpoint, opts)
22   |> Transport.check_origin(handler, endpoint, opts)
23   |> Transport.check_subprotocols(opts[:subprotocols])
24   |> case do
25     %{halted: true} = conn ->
26       {:error, conn}
27
28     %{params: params} = conn ->
29       keys = Keyword.get(opts, :connect_info, [])
30       connect_info = Transport.connect_info(conn, endpoint, keys)
31       config = %{endpoint: endpoint, transport: :websocket, options: opts}
32
33       case handler.connect(config) do
34         {:ok, state} -> {:ok, conn, state}
35         :error -> {:error, Plug.Conn.send_resp(conn, 403, "")}
36         {:error, reason} ->
37           {m, f, args} = opts[:error_handler]
38           {:error, apply(m, f, [conn, reason | args])}
39       end
40   end
41 end
42
43 def connect(conn, _, _, _) do
44   {:error, Plug.Conn.send_resp(conn, 400, "")}
45 end
46
47 def handle_error(conn, _reason), do: Plug.Conn.send_resp(conn, 403, "")
```

The Transport is a functional pipeline that goes into setting up and handling the request

It is NOT a Plug

We can see that handler.connect is called, but our app defines /3 and not /1

Description

[cowboy_websocket](#)

The module `cowboy_websocket` implements Websocket as a Ranch protocol. It also defines a callback interface for handling Websocket connections.

Callbacks

Websocket handlers must implement the following callback interface:

```
init(Req, State)
-> {cowboy_websocket, Req, State}
| {cowboy_websocket, Req, State, Opts}

websocket_init(State)          -> CallResult  %% optional
websocket_handle(InFrame, State) -> CallResult
websocket_info(Info, State)    -> CallResult
```

At the end of the day, Phoenix wraps around Cowboy for its web functionality

Coming Together

- Phoenix uses *cowboy_websocket* to power it's WS implementation
- *Transports.WebSocket* doesn't go through the application router or plug, it's a function pipeline
- Our **AppSocket.connect/1** function is called by *Transports.WebSocket*
- What defines **connect/1**?

Phoenix.Socket connect/1

```
286      @doc false
287      • def connect(map), do: Phoenix.Socket.__connect__(__MODULE__, map, @phoenix_socket_options)
288  end
```

connect /3

```
433      {:ok, serializer} ->
434      • result = user_connect(user_socket, endpoint, transport, serializer, params, connect_info)
435  end
```

user_connect/6

```
552      connect_result =
553      if function_exported?(handler, :connect, 3) do
554        handler.connect(params, socket, connect_info)
555      else
556        handler.connect(params, socket)
557      end
558
559      case connect_result do
560      {:ok, %Socket{} = socket} ->
561        case handler.id(socket) do
562          nil ->
563            {:ok, {state, socket}}
564
565          id when is_binary(id) ->
566            {:ok, {state, %{socket | id: id}}}
```

Phoenix proxies the connection request through multiple layers until it finally gets to our AppSocket.connect function

The result of handler.connect indicates whether the Socket will be allowed, or will fail

How does a Channel Join?

Establish our starting point

— — —

- Network Inspector for WebSocket shows a phx_join being sent to server
- Message format looks foreign (now)

```
↑ ["3","3","room:a","phx_join",{}]
```

```
↓ ["3","3","room:a","phx_reply",{"response":{},"status":"ok"}]
```

Grep for websocket_handle, based on cowboy_websockets interface

```
→ demo grep -R "websocket_handle" deps/**/*.ex
deps/phoenix/lib/phoenix/endpoint/cowboy2_handler.ex: def websocket_handle({opcode, payload}, [handler | state]) when opcode in [:text, :binary] do
deps/phoenix/lib/phoenix/endpoint/cowboy2_handler.ex: def websocket_handle({opcode, payload}, handler_state) when opcode in [:ping, :pong] do
deps/phoenix/lib/phoenix/endpoint/cowboy2_handler.ex: def websocket_handle(opcode, handler_state) when opcode in [:ping, :pong] do
deps/phoenix/lib/phoenix/endpoint/cowboy2_handler.ex: def websocket_handle(_other, handler_state) do
```

Cowboy2Handler.websocket_handle/2

```
---
133   def websocket_handle({opcode, payload}, [handler | state]) when opcode in [:text, :binary] do
134     handle_reply(handler, handler.handle_in({payload, opcode: opcode}, state))
135   end
136
137   def websocket_handle({opcode, payload}, handler_state) when opcode in [:ping, :pong] do
138     handle_control_frame({payload, opcode: opcode}, handler_state)
139   end
140
141   def websocket_handle(opcode, handler_state) when opcode in [:ping, :pong] do
142     handle_control_frame({nil, opcode: opcode}, handler_state)
143   end
144
145   def websocket_handle(_other, handler_state) do
146     {:ok, handler_state}
147   end
```

It's our DemoSocket!

We know that cowboy is used to process websocket messages, so we know that is where we need to look next

Our DemoSocket.handle_in function is called for text or binary messages

Phoenix.Socket handle_in/2

```
292      @doc false
293      def handle_in(message, state), do: Phoenix.Socket.__in__(message, state)
294
```

__in__/2

```
466      def __in__({payload, opts}, {state, socket}) do
467          %{topic: topic} = message = socket.serializer.decode!(payload, opts)
468          handle_in(Map.get(state.channels, topic), message, state, socket)
469      end
```

Wait a second...is this a Channel<->topic mapping?

Socket.handle in phx_join

```
599 defp handle_in(nil, %{event: "phx_join", topic: topic, ref: ref, join_ref: join_ref} = message, state, socket) do
600   case socket.handler.__channel__(topic) do
601     {channel, opts} ->                                     We'll come to this in a minute
602       case Phoenix.Channel.Server.join(socket, channel, message, opts) do
603         {:ok, reply, pid} ->
604           reply = %Reply{join_ref: join_ref, ref: ref, topic: topic, status: :ok, payload: reply}
605           state = put_channel(state, pid, topic, join_ref)
606           {:reply, :ok, encode_reply(socket, reply), {state, socket}}
607
608         {:error, reply} ->
609           reply = %Reply{join_ref: join_ref, ref: ref, topic: topic, status: :error, payload: reply}
610           {:reply, :error, encode_reply(socket, reply), {state, socket}}
611       end
612
613     _ ->
614       {:reply, :error, encode_ignore(socket, message), {state, socket}}
615   end
616 end
```

Note that arg 1 is `nil`. That is the channel_pid and makes sense because the channel hasn't been joined yet

What is Channel.Server? We'll see soon.

Notice lots of error handling. Each layer that we've gone through so far is usually wrapped in error handling

Socket.put_channel

```
650 defp put_channel(state, pid, topic, join_ref) do
651   %{channels: channels, channels_inverse: channels_inverse} = state
652   monitor_ref = Process.monitor(pid)
653
654   %{
655     state |
656     channels: Map.put(channels, topic, {pid, monitor_ref}),
657     channels_inverse: Map.put(channels_inverse, pid, {topic, join_ref})
658   }
659 end
---
```

Understanding topics and Channels is one of the most core pieces of real-time phoenix

A topic is a key in a map. Used for routing

A Channel is a process that responds to messages from the client

Regroup

- Cowboy handler routes all incoming messages through **DemoSocket**
- Each Socket is a new Process
- Top-level functions handled by **use Phoenix.Socket**
- Channels are just an entry in a map
- Each Channel must be a process (can be monitored)

Back to Our Questions

- What is a Socket, and how does it connect?
- How does a Channel differ from a Socket?
- How does a Channel become joined?
- How does a message get processed and responded to?

Channel Startup

Channel.Server join/4

```
18 def join(socket, channel, message, opts) do
19   %{topic: topic, payload: payload, ref: join_ref} = message
20   assigns = Map.merge(socket.assigns, Keyword.get(opts, :assigns, %{}))
21   socket = %{socket | topic: topic, channel: channel, join_ref: join_ref, assigns: assigns}
22
23   ref = make_ref()
24   from = {self(), ref}
25   child_spec = channel.child_spec({socket.endpoint, from})
26
27   case PoolSupervisor.start_child(socket.endpoint, socket.handler, from, child_spec) do
28     {:ok, pid} ->
29       send(pid, {Phoenix.Channel, payload, from, socket})
30       mon_ref = Process.monitor(pid)
31
32       receive do
33         {^ref, {:ok, reply}} ->
34           Process.demonitor(mon_ref, [:flush])
35           {:ok, reply, pid}
36
```

A GenServer is started and sent a message before the Channel is joined

A Supervisor creates a child process for a Channel.Server

The process is created BEFORE it's attempted to be joined. This allows you to send(self()) messages in the join/3 function and it works as expected

Channel.Server handle_info initial

```
287 def handle_info({Phoenix.Channel, auth_payload, {pid, _} = from, socket}, ref) do
288   Process.demonitor(ref)
289   %{channel: channel, topic: topic, private: private} = socket
290   Process.put(:"$callers", [pid])
291
292   socket = %{
293     socket
294     | channel_pid: self(),
295     private: Map.merge(channel.__socket__(:private), private)
296   }
297
298   start = System.monotonic_time()
299   {reply, state} = channel_join(channel, topic, auth_payload, socket)
300   duration = System.monotonic_time() - start
301   metadata = %{params: auth_payload, socket: socket, result: elem(reply, 0)}
302   :telemetry.execute([:phoenix, :channel_joined], %{duration: duration}, metadata)
303   GenServer.reply(from, reply)
304   state
305 end
```

This function processes the initial request to join a Channel

Line 299 is invocation of the join function, everything else is setup or monitoring code

Channel.Server channel_join

```
376 defp channel_join(channel, topic, auth_payload, socket) do
377   case channel.join(topic, auth_payload, socket) do
378     {:ok, socket} ->
379       {{:ok, %{}}, init_join(socket, channel, topic)}
380
381     {:ok, _} ->
```

init_join

```
400 defp init_join(socket, channel, topic) do
401   %{transport_pid: transport_pid, serializer: serializer, pubsub_server: pubsub_server} = socket
402   ...
418 Process.monitor(transport_pid)
419 fastlane = {:%fastlane, transport_pid, serializer, channel.__intercepts__()}
420 PubSub.subscribe(pubsub_server, topic, metadata: fastlane)
421
422   {:noreply, %{socket | joined: true}}
423 end
```

Hey, it's the PubSub
subscription!

If you feel curious, you can dig into “fastlane” here. I typically would because it would be an unknown if it's the first time I've seen it

Channel.Server handle_info_proxy

```
341 def handle_info(msg, %{channel: channel} = socket) do
342   if function_exported?(channel, :handle_info, 2) do
343     msg
344     |> socket.channel.handle_info(socket)
345     |> handle_result(:handle_info)
346   else
347     warn_unexpected_msg(:handle_info, 2, msg)
348     {:noreply, socket}
349   end
350 end
```

Our app's Channel is not actually a GenServer module, it's a proxy for one!

Very important to understand the Channel process structure. What are the processes (Channel.Server, cowboy_websocket handler), process facades (YourChannel), or handlers (YourSocket)?

Regroup: What is a Channel?

- GenServer backed by Phoenix.Channel.Server
- A Channel can receive any unhandled messages via `handle_info`
- Lives under a Phoenix.Socket, mapped by topic

Channel Incoming Message & Reply

Phoenix.Socket handle_in

```
---
629 defp handle_in({pid, _ref}, message, state, socket) do
630   send(pid, message)
631   {:ok, {state, socket}}
632 end
633
```

This looks pretty simple, it just sends a message to the channel process

What happens when a Socket gets a new message? It calls this function (or another, based on pattern-matching)

[Channel.Server handle_info](#)

```
---
311 def handle_info(
312     %Message{topic: topic, event: event, payload: payload, ref: ref},
313     %{topic: topic} = socket
314 ) do
315     start = System.monotonic_time()
316     result = socket.channel.handle_in(event, payload, put_in(socket.ref, ref))
317     duration = System.monotonic_time() - start
318     metadata = %{ref: ref, event: event, params: payload, socket: socket}
319     :telemetry.execute([:phoenix, :channel_handled_in], %{duration: duration}, metadata)
320     handle_in(result)
321 end
```

Again, pattern matching plays a big part here. There are other “admin” type messages to handle

A lot of code that we see in these files is a single important line surrounded by monitoring or error handling

Channel.Server handle_in result

Replies are handled
based on the Channel
response

```
487     ## Handle in/replies
488
489     defp handle_in({:reply, reply, %Socket{} = socket}) do
490       handle_reply(socket, reply)
491       {:noreply, put_in(socket.ref, nil)}
492     end
493
494     defp handle_in({:stop, reason, reply, socket}) do
495       handle_reply(socket, reply)
496       handle_result({:stop, reason, socket}, :handle_in)
497     end
498
499     defp handle_in(other) do
500       handle_result(other, :handle_in)
501     end
502
```

Pattern matching is leveraged to decide how to handle different reply formats

Channel.Server handle_reply

```
503     defp handle_reply(socket, {status, payload}) when is_atom(status) and is_map(payload) do
504         reply(
505             socket.transport_pid,
506             socket.join_ref,
507             socket.ref,
508             socket.topic,
509             {status, payload},
510             socket.serializer
511         )
512     end
```

transport_pid = cowboy_websocket handler process

reply

```
248     def reply(pid, join_ref, ref, topic, {status, payload}, serializer)
249         when is_binary(topic) and is_map(payload) do
250         reply = %Reply{topic: topic, join_ref: join_ref, ref: ref, status: status, payload: payload}
251         send(pid, serializer.encode!(reply))
252         :ok
253     end
254
```

A reply is a message being sent to a process (cowboy_websocket transport)

All messages have a particular format dictated by the Serializer

Socket.V2.JSONSerializer.encode!

```
14 def encode!(%Reply{} = reply) do
15   data = [
16     reply.join_ref,
17     reply.ref,           Messages are converted to/from this
18     reply.topic,         array format
19     "phx_reply",
20     %{status: reply.status, response: reply.payload}
21   ]
22
23   {:socket_push, :text, Phoenix.json_library().encode_to_iodata!(data)}
24 end
--
```

Regroup: How are Channel Messages Processed?

- Each message is handled by the Channel Server
- A given message is processed serially, one at a time
- The Socket transport is sent the serialized reply
- All messages over the wire (up/down) are serialized into a particular format

Back to Our Questions

- What is a Socket, and how does it connect?
- How does a Channel differ from a Socket?
- How does a Channel become joined?
- How does a message get processed and responded to?
- We have (rough) answers to everything!

Quick Dive Into LiveView (time permitting)

- Feels more familiar than unknown
- LiveView mounts a Phoenix.Socket module called **Phoenix.LiveView.Socket**
- **Phoenix.LiveView.Channel** follows the Phoenix.Channel behaviour contract
- It implements all of the important bits, but doesn't *use* *Phoenix.Channel*

**The Phoenix real-time stack is incredibly
powerful.**

**You write not a lot of application code, but get
rich real-time features.**

**Read through the source code if you hit a snag,
or are just curious.**

You may unblock yourself or learn something new, but you'll definitely become more self-sufficient.

Thank You!

- Book: sal.es/rtp
- Twitter: @yoooooaaaaa
- Email: steve@salesloft.com
- Slides with Notes: <https://bit.ly/anatomy-real-time-elixir>